
Workstream: Observability, CI/CD & Testing Pipeline

Period covered: Week 3

Date: 18 Sep 2026

Source: docs/week-3-report.docx (Week 3 Task Breakdown — RESOLV-HQ), Tasks 18–20.

Critical path note: "The critical path for Week 3 is the same one flagged at the end of Week 2 — live model + retrieval wiring (Tasks 13-15, 17) — plus everything downstream of it that currently has nothing to test against (Tasks 18-20)." All three of this workstream's tasks fall in that "nothing to test against yet" category.

Tasks

Task #	Task Description	Notes
18	Centralized Audit Logging Pipeline — Set up structured JSON log aggregators to capture agent intermediate reasoning, tool inputs/outputs, system timestamps, and session context.	admin_activity_logs already captures structured JSON detail for two call sites — approval_decided (lib/approval-decisions.ts) and request_assigned (routes/admin/approvals.ts) — with admin id, action, target type/id, and timestamp. It is not yet a general-purpose pipeline: there's no logging of agent intermediate reasoning or tool inputs/outputs, since no live agent/tool-calling loop exists yet (depends on Tasks 14/15/17 landing first).
19	GitHub Actions Integration Test Pipeline — Configure automated CI workflows to validate Docker container builds, lint codebases, and run endpoint sanity checks on pull requests.	The only content under .github/ across the three repos is an unrelated java-upgrade modernization scaffold — no workflow YAML exists for Docker build validation, linting, or endpoint sanity checks on PRs.
20	End-to-End API Integration Test Suite — Build automated Postman/Pytest integration suites to simulate full user request journeys through gateway, model, tool, and approval stages.	No test files (*.test.ts, *.spec.ts) or Postman collections exist in resolv-hq, resolv-hq-backend, or resolv-hq-customer today. This depends on Tasks 7, 14, 15, and 17 landing first so there's an actual end-to-end journey (gateway → model → tool → approval) to script against.

Summary

This workstream is the most blocked of the six: all three tasks depend on work landing elsewhere first (the live model/retrieval pipeline, the SSE streaming pipeline, and the post-approval dispatch engine) before there's a real journey to log, test, or validate in CI. The one piece already in place — structured activity logging for approval decisions and assignments — is a solid foundation to extend once the agent loop is live.

Plan for next week

- Track progress on Tasks 7, 13–15, and 17 closely — this workstream can't meaningfully start Tasks 19–20 until at least one real gateway → model → tool → approval journey exists.
- Extend `admin_activity_logs` (or a purpose-built log table) to capture agent intermediate reasoning and tool inputs/outputs once the agent loop is live (Task 18).
- Stand up a baseline GitHub Actions workflow now for what's already testable — Docker build validation and linting — rather than waiting on the full pipeline (Task 19).
- Draft the E2E test suite's scenario list now against the target journey, so it's ready to implement the moment Task 17 lands (Task 20).